

Deep Learning (TensorFlow, Keras) with ResNet50: Image Binary Classifier (Part 2)

In this project, a model is trained to perform binary classification for cats and dogs pictures. The pretrained model ResNet50 is used. This document is the second part of the whole training process.

Iteration 2: Model retraining with data augmentation, fine-tuning (last 10 layers only) and learning_rate = 1e-5 (smaller)

```
# (height, width, channels)
input_shape = (224, 224, 3)
batch_size = 8
learning_rate = 1e-5
neurons = 128
path_dataset = '../dataset_cat_dogs'
folder_cat = 'Cat'
folder_dog = 'Dog'
folder_models = '../models'

import pandas as pd
import matplotlib.pyplot as plt
import os
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping,
ReduceLROnPlateau, ModelCheckpoint
from tensorflow.keras.models import load_model

# Find how many cats and dogs images exist
cat_imgs = os.listdir(os.path.join(path_dataset, folder_cat))
dog_imgs = os.listdir(os.path.join(path_dataset, folder_dog))
print(f'Cat images found: {len(cat_imgs)}')
print(f'Dog images found: {len(dog_imgs)}')
```

Classes are balanced.

Data augmentation

```
def load_data(path, input_shape=input_shape, batch_size=batch_size,
seed=123, validation_split=0.2):
    """Function to create 2 ImageDataGenerators to split dataset into
    train and validation datasets.
    Data augmentation is not implemented for the validation
    dataset."""
    height, width = input_shape[:2]
    datagen = ImageDataGenerator(rescale=1.0/255, zoom_range=0.15,
        horizontal_flip=True, vertical_flip=False,
        height_shift_range=0.15, width_shift_range=0.15,
        brightness_range=(0.8, 1.2), rotation_range=20,
        validation_split=validation_split
    )
    train_data = datagen.flow_from_directory(path,
        target_size=(height, width), batch_size=batch_size,
        class_mode='binary', subset='training', seed=seed
    )
    val_datagen = ImageDataGenerator(rescale=1.0/255,
        validation_split=validation_split
    )
    val_data = val_datagen.flow_from_directory(path,
        target_size=(height, width), batch_size=batch_size,
        class_mode='binary', subset='validation', seed=seed
    )
    return train_data, val_data

# Split training and validation datasets
train, val = load_data(path_dataset)

print(f"Classes found: {train.class_indices}")
print(f"Training images: {train.samples}")
print(f"Validation images: {val.samples}")

# Obtain images and target
images, labels = next(train)

# Show 8 training images (batch_size=8)
figure, axes = plt.subplots(nrows=2,ncols=4, figsize=(8, 6))
for item in zip(axes.ravel(), images, labels):
    axes, image, target = item
    axes.imshow(image)
    axes.set_title(f'Target: {target:.0f}')
    axes.set_xticks([])
    axes.set_yticks([])
plt.tight_layout()
plt.show()

# Image dimensions
print(images.shape)
```

Model retraining (iteration 2)

```
def train_model(model, train_data, val_data, epochs, version_model,
                folder_models=folder_models):
    """Function to train the model and save the best one
    according to the validation accuracy."""
    file_name =
os.path.join(folder_models, f'binary_model_v{version_model}.h5')
    callbacks = [
        EarlyStopping(monitor='val_loss', patience=5,
            restore_best_weights=True, verbose=0),
        ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=3,
            min_lr=1e-6, verbose=0),
        ModelCheckpoint(file_name, monitor='val_accuracy',
            save_best_only=True, verbose=1)
    ]

    history = model.fit(train_data, validation_data=val_data,
                        epochs=epochs, callbacks=callbacks, verbose=2)

    return model, history

# Load model v1
model_v2 =
load_model(os.path.join(folder_models, 'binary_model_v1.h5'))
model_v2.summary()

epochs = 20
version_model = 2
print(f"Parameters: batch_size = {batch_size}, learning_rate =
{learning_rate}, neurons = {neurons}, epochs = {epochs}")

# last 10 layers
for layer in model_v2.layers[0].layers[-10:]:
    layer.trainable = True

# Recompile
model_v2.compile(optimizer=Adam(learning_rate=learning_rate),
                 loss='binary_crossentropy', metrics=['accuracy'])

print(f"TensorFlow Version: {tf.__version__}")

# Ensure GPU is available
physical_devices = tf.config.list_physical_devices('GPU')
if len(physical_devices) > 0:
    tf.config.experimental.set_memory_growth(physical_devices[0], True)
    print("GPU is available and memory growth is enabled.")
else:
    print("GPU not available, training will be on CPU.")

# Retrain the model
```

```
model_v2, history_stage2 = train_model(model_v2, train, val,  
epochs=epochs, version_model=version_model)
```

Result 2: val_accuracy=?%.

```
pd.DataFrame(history_stage2.history).plot(figsize=(12, 4))  
plt.show()  
  
# Save model  
#  
model.save(os.path.join(folder_models, f'binary_model_v{version_model}.  
keras'))
```

In the next iteration, the model will be retrained, data augmentation and fine-tuning (the last 20 layers) will be performed.